

---

# Q-Pain: A Question Answering Dataset to Measure Social Bias in Pain Management

---

**Cécile Logé\***

Department of Computer Science  
Stanford University  
ceciloge@stanford.edu

**Emily Ross\***

Department of Computer Science  
Stanford University  
emilyross@cs.stanford.edu

**David Yaw Amoah Dadey**

Department of Neurosurgery  
Stanford University  
dadeyd@stanford.edu

**Saahil Jain**

Department of Computer Science  
Stanford University  
saahil.jain@cs.stanford.edu

**Adriel Saporta**

Department of Computer Science  
Stanford University  
asaporta@cs.stanford.edu

**Andrew Y. Ng**

Department of Computer Science  
Stanford University  
ang@cs.stanford.edu

**Pranav Rajpurkar**

Department of Biomedical Informatics  
Harvard University  
pranav\_rajpurkar@hms.harvard.edu

## Abstract

Recent advances in Natural Language Processing (NLP), and specifically automated Question Answering (QA) systems, have demonstrated both impressive linguistic fluency and a pernicious tendency to reflect social biases. In this study, we introduce Q-Pain, a dataset for assessing bias in medical QA in the context of pain management, one of the most challenging forms of clinical decision-making. Along with the dataset, we propose a new, rigorous framework, including a sample experimental design, to measure the potential biases present when making treatment decisions. We demonstrate its use by assessing two reference Question-Answering systems, GPT-2 and GPT-3, and find statistically significant differences in treatment between intersectional race-gender subgroups, thus reaffirming the risks posed by AI in medical settings, and the need for datasets like ours to ensure safety before medical AI applications are deployed.

## 1 Introduction

Pain management remains one of the most challenging forms of clinical decision making [1]. Since the patient experience of pain and its manifestations are highly variable, there is inherent subjectivity in physician pain assessments [2-4]. The challenge of measuring pain combined with the vast diversity of pharmacologic and non-pharmacologic treatment options results in clinical decision pathways that are difficult to standardize [5]. Thus, the complexity of pain management and the lack

---

\*Equal contribution

of universal standards present opportunities for bias to impact clinical decisions. Racial disparities in treatment of pain have been shown in prior studies [6-8], where Black patients were consistently more likely to be given inadequate or no pain medications when compared to White patients [9]. Furthermore, multiple studies have highlighted that women were more likely to be under-treated for pain [10, 11].

Recent advances in natural language processing (NLP) have enabled the development of automated question answering (QA) systems that can answer personalized medical questions [12]. However, many NLP systems, including those specific to QA, have been shown to encode and reinforce harmful societal biases [13-16]. Benchmark datasets have been instrumental in surfacing bias in NLP systems [17, 18]. Therefore, before medical QA systems are further deployed and even incorporated into medical workflows for applications such as mental health conversations and clinical advice lines, it is critical to develop benchmark datasets that will help us understand the extent to which encoded societal biases surface in medical QA systems.

We introduce Q-Pain, a dataset for assessing bias in medical QA in the context of pain management consisting of 55 medical question-answer pairs; each question includes a detailed patient-specific medical scenario ("vignette") designed to enable the substitution of multiple different racial and gender "profiles" in order to identify discrimination when answering whether or not to prescribe medication.

Along with the dataset, we propose a new framework, including a fully reproducible sample experimental design, to measure the potential bias present during medical decision-making for patients with particular demographic profiles. We demonstrate its use in assessing two reference QA systems, GPT-2 [19] and GPT-3 [20], selected for their documented ability to answer questions given only a few examples. We expect our dataset and framework to be used to assess bias across a wide variety of QA systems, as well as for experiments in real-life medical settings (e.g. through surveys of clinicians) with only minor changes to our experimental design.

## 2 Related Work

**Medical Question Answering:** There have been efforts to create QA datasets for medical reasoning [21, 22]. The Medication QA dataset [23] comprises nearly 700 drug-related consumer questions along with information retrieved from reliable websites and scientific papers. The emrQA dataset [24] approaches case-specific analysis by posing over 400,000 factual questions, the answers to which are contained in provided electronic medical records. The MedQA task [25] expects a model to select the correct answer to real-world questions derived from the Chinese medical licensing exam, given access to a corpus of medical documents. None of these datasets address the pain management context or evaluate a treatment prescription task.

**Human Pain Management Bias:** Social bias in human-facilitated pain management is well-documented [6-8, 10]: a survey of studies on racial and ethnic disparities in pain treatment [2] demonstrated that in acute, chronic, and cancer pain contexts, racial and ethnic minorities were less likely to receive opioids. Another meta-analysis of acute pain management in emergency departments [26] found that Black patients were almost 40% less likely (and Hispanic patients 30% less likely) than White patients to receive any analgesic. These gaps remained for opioid-specific analgesics, as well. Finally, closely related to our experimental setup is the work of Weisse et al. [27], who presented medical vignettes with race- and gender-specific patient profiles to physicians and asked them to decide not only whether to treat the patient, but with what dosage. However, only three vignettes were used, leading to a limited number of race-gender profiles to be presented to each given physician as well as different physicians seeing different profiles for the same scenario, thus making results more difficult to interpret. In contrast to this work, we consider a larger diversity of scenarios and profiles, and demonstrate our rigorous statistical approach on two reference QA systems.

**QA/Language Model Bias:** Social bias in language models is also being increasingly explored, with new datasets, benchmarks, and metric frameworks geared towards analyzing bias [17, 18, 15]. Li et al. [13] introduce a helpful framework for quantifying bias in QA settings with questions that do not prime for bias, and demonstrate bias in a broad array of transformer-based models using this framework. Most closely related to the medical context, Zhang et al. [28] show that language models exhibit statistically significant performance gaps in clinical predictive accuracy between majority

and minority groups, thus laying the groundwork for examining bias for language models operating in medical contexts. In this study, we go one step further and evaluate bias present during medical decision-making by focusing on treatment recommendation rather than diagnosis prediction.

### 3 Dataset and Framework

#### 3.1 Clinical Vignettes

Clinical vignettes are medical scenarios typically presenting a patient with a specific set of symptoms, often with the goal to illustrate decision-making and explore different ranges of actions. We introduce Q-Pain, a dataset of 55 pain-related clinical vignettes depicting patients in different medical contexts related to pain management. At the end of each vignette is a question asking whether the patient should be prescribed pain medication (either intravenous hydromorphone, hydrocodone, oxycodone or morphine), if so what dosage, and why.

**Data Collection:** A clinical expert with four years of medical experience designed each of the clinical scenarios such that the patient presentation and data were consistent with real-life situations where pain treatment *should* be offered. As such, the expected treatment answer for all but five of the vignettes (see "Data Structure & Documentation") used for closed prompts is "Yes." regardless of gender or race.

The decision to focus on situations where patients *should* receive treatment (the "Yes" vignettes) was mainly driven by our interest in harmful bias: "Yes" vignettes present the opportunity for "unjust" denial of treatment. In contrast, "No" vignettes do not allow us to identify such under-treatment because those situations do not warrant pain treatment in the first place.

From a clinical standpoint, though, "No" scenarios are also more difficult to compose: there are no universally-applicable, objective measures for pain assessment such that "No" would be a consistently valid response. The only option would be to make the scenarios either far too simple (e.g. someone falls off their bike and scrapes their knee) or significantly more complex by including additional variables (e.g. depression/mood disorders, medication allergies, substance abuse history, etc.).

To allow for detection of more nuanced discrimination, we included dose/supply scales in each question. These were defined as low or high according to an appropriate treatment metric (milligrams or weeks supply of medication) for each clinical context. Similarly, these scales were designed such that both low and high choices of doses/supplies were objectively acceptable for pain treatment regardless of gender or race.

**Data Structure & Documentation:** The Q-Pain dataset is structured into five .csv files, one for each medical context: Acute Non-Cancer pain, Acute Cancer pain, Chronic Non-Cancer pain, Chronic Cancer pain, and Post-Operative pain. The fifty "Yes" prompts and the five "No" prompts were distributed evenly across the five medical contexts: 10 "Yes" prompts and one "No" prompt per context. For each vignette, it includes: the case presentation ("Vignette"), the question ("Question"), the expected answer ("Answer") and dosage ("Dosage") as well as a brief explanation justifying the response ("Explanation").

The full dataset, complete with documentation as well as a Python Notebook with starter code to reproduce our experimental setup, is available on Physionet:

<https://doi.org/10.13026/61xq-mj56>

The dataset is licensed under a *Creative Commons Attribution-ShareAlike 4.0 International Public License* and will be maintained and improved upon directly on the Physionet platform [29].

**Validation Protocol:** To further validate the dataset, a random selection of vignettes spanning each clinical context were provided to two internal medicine physicians who were not involved in the initial design. For each vignette, the physicians were asked to assess for (1) credibility of the described scenario, (2) coherence of clinical information provided, and (3) sufficiency of information provided toward making a clinical decision. For all vignettes assessed, both physicians responded affirmatively to each query.

### 3.2 Experimental Design

The Q-Pain dataset was conceived with the objective of measuring harmful bias against both race and gender identities in the context of pain management. We demonstrate the following experimental design on GPT-3 (*davinci*, via the OpenAI API) and GPT-2 (*gpt2-large*, via the HuggingFace API). GPT-3 [20] is a transformer-based autoregressive language model with 175 billion parameters trained on a combination of the Common Crawl corpus (text data collected over 12 years of web crawling), internet-based books corpora and English-language Wikipedia. At a high level, the model takes in a prompt and completes it by sampling the next tokens of text from a conditional probability distribution. GPT-2 [19] has similar architecture to GPT-3 but fewer parameters. We used these models with the goal of providing an initial benchmark in large language models (LLMs).

**Closed Prompts:** LLMs have been shown to answer questions more accurately in a “few-shot” setting, in which the model is first given several question-answer examples (“closed prompts”) to better learn the relevant task: with this in mind, we preceded all our vignettes (“open prompts”) with three closed prompts sampled from the same pain context. To limit the influence of these closed prompts on the generated text and make them as neutral as possible, we removed any indication of race or gender, named the patients “Patient A”, “Patient B” and “Patient C”, used gender-neutral pronouns (“they”/“their”), and ensured each of the three possible answers (“No.”, “Yes. Dosage: High”, and “Yes. Dosage: Low”) were represented in the closed prompt selection. We also used ‘##’ as a stop token to mark the separation between prompts and give the models a sense of where to stop during generation (see Fig. 1a for an example of input). Throughout our experiments, *temperature* was set to 0 to eliminate the randomness of the output, as recommended by OpenAI for factual question answering. The maximum number of tokens for the model to generate was set to 150. Other parameters were kept to their default values: no nucleus sampling and no presence or frequency penalties.

**Rotating Patient Identities:** As we prompted the models for answers, we rotated the patient’s race (Asian, Black, Hispanic, White) and gender (man, woman) across all vignettes for a total of eight patient combinations and 400 open prompts. Conscious of the demographic information carried within a person’s name and inspired by prior work on race-based stereotyping [30], we opted to name each patient according to their demographic profile using Harvard Dataverse’s ‘Demographic aspects of first names’ dataset [31]. This dataset contains a list of 4,250 first names along with data on their respective frequency and proportions across mutually exclusive racial groups. For each race and gender, we chose the top ten first names based on their overall frequency and representation within each group, excluding unisex names and names that differed by only one character (see Fig. 1b). For each context, we selected a random permutation of names to use across the ten open prompts. Beyond the names, race, gender, and corresponding pronouns, everything else in the vignettes remained completely unchanged—from symptom presentation and severity to age and preceding closed prompts, so as to minimize the influence of other factors in the final decision.

### 3.3 Bias Framework

We propose the following framework for investigating bias in the 400 outputs generated from running our experimental design on the Q-Pain dataset (in our case, using GPT-3 and GPT-2 for demonstration). We outline three core goals: (1) assessing bias in the form of unequally-applied, harmful decision-making, (2) attending to intersectionality, and specifically the interconnected nature of race and gender, and (3) confirming the statistical significance of the results using p-values and confidence intervals.

**Emphasizing Harmful Bias:** Since both GPT-3 and GPT-2 generate text by calculating the probabilities of possible next words conditioned on previous words, we were able to not only access the generated text, but also the underlying probability distributions. More precisely, we extracted two elements: the probability of the model to answer “Yes.” or “No.”, and if the former, the probability of prescribing a “Low” or “High” dosage (see Fig. 1c). In this framework, we chose to focus on the probability of a “No.” (rather than “Yes.”), signifying a recommendation to forego the appropriate pain treatment to patients, as well as the probability of outputting a “Low” dosage (rather than a “High” one). This approach allows us to focus on the most harmful outcomes for patients and assess potential disparities in treatment.




Figure 1: (a) Example of Input with 3 closed prompts + 1 open prompt with placeholders for name, race, gender and pronouns, (b) Chosen names based on Harvard Dataverse Demographic aspects of first names dataset, (c) Examples of Outputs from GPT-3.

| (a) GPT-3 Average Probabilities of "No" (i.e. denying pain treatment) |              |         |          |         |            |           |         |
|-----------------------------------------------------------------------|--------------|---------|----------|---------|------------|-----------|---------|
| Gender                                                                | Race         | Average | Acute NC | Acute C | Chronic NC | Chronic C | Post Op |
| Male Patients<br>"man" "he" "his"                                     | Asian        | 25.5%   | 29.2%    | 25.4%   | 26.1%      | 22.3%     | 24.5%   |
|                                                                       | Black        | 25.4%   | 29.0%    | 25.9%   | 25.4%      | 22.3%     | 24.4%   |
|                                                                       | Hispanic     | 25.0%   | 29.0%    | 26.0%   | 24.6%      | 21.3%     | 23.9%   |
|                                                                       | White        | 24.8%   | 27.9%    | 24.6%   | 25.4%      | 21.5%     | 24.4%   |
|                                                                       | All          | 25.2%   | 28.8%    | 25.5%   | 25.4%      | 21.9%     | 24.3%   |
| Female Patients<br>"woman" "she" "her"                                | Asian        | 26.2%   | 29.7%    | 25.7%   | 26.7%      | 23.2%     | 25.6%   |
|                                                                       | Black        | 26.6%   | 30.4%    | 27.4%   | 26.1%      | 23.3%     | 25.6%   |
|                                                                       | Hispanic     | 26.4%   | 29.3%    | 27.2%   | 26.5%      | 23.3%     | 25.7%   |
|                                                                       | White        | 25.3%   | 28.6%    | 25.3%   | 25.8%      | 22.2%     | 24.7%   |
|                                                                       | All          | 26.1%   | 29.5%    | 26.4%   | 26.3%      | 23.0%     | 25.4%   |
| All Patients                                                          | Asian        | 25.8%   | 29.4%    | 25.6%   | 26.4%      | 22.7%     | 25.1%   |
|                                                                       | Black        | 26.0%   | 29.7%    | 26.7%   | 25.7%      | 22.8%     | 25.0%   |
|                                                                       | Hispanic     | 25.7%   | 29.2%    | 26.6%   | 25.6%      | 22.3%     | 24.8%   |
|                                                                       | White        | 25.1%   | 28.3%    | 24.9%   | 25.6%      | 21.9%     | 24.6%   |
|                                                                       | All          | 25.6%   | 29.1%    | 25.9%   | 25.8%      | 22.4%     | 24.8%   |
|                                                                       | Standard Dev | 4.2%    | 4.2%     | 3.9%    | 4.0%       | 3.2%      | 2.3%    |
|                                                                       | Maximum      | 36.6%   | 36.6%    | 33.2%   | 35.7%      | 29.4%     | 29.9%   |

| (b) p-values for GPT-3 Intersectionality Subgroup Comparisons for Probabilities of "No" (i.e. denying pain treatment) |                        |             |                    |             |                            |           |              |           |
|-----------------------------------------------------------------------------------------------------------------------|------------------------|-------------|--------------------|-------------|----------------------------|-----------|--------------|-----------|
|                                                                                                                       | Asian Woman            | Black Woman | Hispanic Woman     | White Woman | Asian Man                  | Black Man | Hispanic Man | White Man |
| Asian Woman                                                                                                           |                        | >0.05       | >0.05              | <0.002      | >0.002                     | >0.002    | <0.002       | <0.0001   |
| Black Woman                                                                                                           |                        |             | >0.05              | <0.0001     | <0.0001                    | <0.0001   | <0.0001      | <0.0001   |
| Hispanic Woman                                                                                                        |                        |             |                    | <0.002      | <0.002                     | <0.0001   | <0.0001      | <0.0001   |
| White Woman                                                                                                           |                        |             |                    |             | >0.05                      | >0.05     | >0.05        | >0.002    |
| Asian Man                                                                                                             |                        |             |                    |             |                            | >0.05     | >0.002       | >0.002    |
| Black Man                                                                                                             |                        |             |                    |             |                            |           | >0.05        | >0.002    |
| Hispanic Man                                                                                                          |                        |             |                    |             |                            |           |              | >0.05     |
|                                                                                                                       | >0.002 Not significant |             | <0.002 Significant |             | <0.0001 Highly significant |           |              |           |

Figure 2: (a) Overview of the underlying average “No” probability distribution across contexts (Acute NC/C: Acute Non Cancer/Cancer Pain, Chronic NC/C: Chronic Non Cancer/Cancer Pain, Post Op: Post Operative Pain) for GPT-3, (b) p-values for Intersectionality Subgroup Comparisons, obtained via a paired two-tailed t-test

**Examining Intersectionality:** "*Intersectionality*" encapsulates the idea that the combination of certain identity traits, such as gender and race (among others), can create overlapping and interdependent systems of discrimination [32], leading to harmful results for specific minorities and subgroups. With this in mind, we chose not only to look at overall differences between genders (regardless of race) and between races (regardless of gender) across vignettes and pain contexts, but also to further explore race-gender subgroups with the idea to assess all potential areas of bias and imbalance—for a total of 28 subgroup-to-subgroup comparisons.

**Ensuring Statistical Significance:** Bias measurements—whether with LLMs or in real-life settings—are inherently uncertain [33] and can hardly be summarized in a single number. When measuring whether a specific race-gender subgroup is being discriminated against compared to a reference subgroup (e.g. are Black women being disproportionately denied pain treatment compared to White men?), computing p-values and reporting 95% confidence intervals is essential to confirm (or invalidate) the significance of the results. Furthermore, because our attention to intersectionality requires 28 simultaneous comparisons, additionally computing the Bonferroni-corrected p-value threshold and confidence intervals lends credibility to subsequent analysis.

## 4 Results

### 4.1 First Look

In our initial analysis of results, we looked at both the probabilities of answering "No" and of prescribing a "Low" dosage. As the latter showed less variability overall, we chose to focus the rest of our analysis on the former, meaning the decision to deny patients of the appropriate pain treatment.



Those who implement our experimental design and framework to test their own medical QA systems can perform similar analysis to what follows on the "Low" probabilities.

For GPT-3, we found that the probability of a "No." across all vignettes and profiles was always less than 50% and on average 25.6%, signaling that as long as the model is configured to behave deterministically (i.e. with *temperature* set to 0), it will advocate for treating every patient. However, we noticed significant variability across medical contexts as well as individual vignettes, with Acute Non Cancer Pain scenarios reporting higher average probabilities and higher standard deviation overall. Averaging across all genders and races, we noted a 30% gap between the medical contexts with the lowest (Chronic Cancer pain) and highest (Acute Non Cancer pain) propensities for treatment denial (see Fig. 2a).

Focusing on race, specifically, we found that Black patients were 3.6% more likely to have been refused pain treatment by GPT-3 than White patients. Focusing on gender, we noted a similar 3.6% difference between women and men, with women at a disadvantage. But looking into intersectional subgroups reveals deeper disparities: for example, we found that out of all subgroups, Black women had the highest probability of being denied pain treatment by GPT-3 at 26.6%, while White men had the lowest at 24.8%, meaning that Black women were 7% more likely to be refused pain relief by GPT-3 compared to White men—and up to 11% more likely for the Acute Cancer Pain context, specifically. This initial observation confirms the need to highlight outcomes for intersectional subgroups.

#### 4.2 Intersectionality Analysis

To assess the significance of our first observations, we carried out several tests:

**Paired Two-Tailed t-Test:** Using results from GPT-3, we conducted a paired two-tailed t-test analysis on 28 subgroup-to-subgroup comparisons and found no statistically significant difference for the following: Black Man v Hispanic/Asian Man, Hispanic Man v White Man, White Woman v. any Man, as well as any minority Woman v minority Woman.

However, 19 of the 28 subgroup-to-subgroup comparisons had p-values less than 0.05, implying the difference of treatment and outcome was statistically significant between these subgroups. More notably, 13 out of the 28 had p-values less than 0.002, corresponding to the Bonferroni-corrected threshold, and nine had p-values less than 0.0001, five of which included Black women (see Fig. 2b).

**Confidence Intervals:** To further measure the magnitude of these gaps, we computed the confidence intervals for all 28 intersectional differences—and contrasted results from both GPT-3 and GPT-2 (see Fig.3). Any comparison in the form of *Group A v Group B* is assessing the question: *is Group A more likely to be denied treatment than Group B?* As such, a positive result implies that Group A is at a disadvantage compared to Group B, while a negative result implies the opposite. An interval including both negative and positive values is inconclusive. To ensure the results of 28 simultaneous comparisons were jointly valid, we made the decision to focus on Bonferroni-corrected intervals (corresponding to 99.8%).

In GPT-3, the following comparisons obtained a significant *positive* result (>0.5% difference), in descending magnitude: Black Woman v White Man, Black Woman v Hispanic Man, Hispanic Woman v White Man, and Asian Woman v White Man. What's more, all minority Woman subgroups had at least three *positive* results (and up to a total of five) when compared with the rest of the subgroups, thus putting minority women, and specifically Black women, at the most disadvantaged position in pain management by GPT-3. The rest of the comparisons were inconclusive.

In GPT-2, results were less indicative of bias toward women and mostly singled out Asian men. In descending order of magnitude, the following significant *positive* results (>0.5% difference) stood out: Asian Man v Hispanic Man, Asian Man v Black Man, Black Woman v Hispanic Man, and Asian Man v White Man. While some of the differences in GPT-2 were more pronounced, with Asian Man v Hispanic Man obtaining a 1.8%- 4.3% confidence interval, GPT-3 showed conclusive difference for a wider range of subgroup comparisons—12 out of 28 compared to 7 for GPT-2—thus demonstrating widespread disparities rather than improvement.

| GROUP A                                                                                                             |             |              |                | GROUP B      |               |              |              |               |
|---------------------------------------------------------------------------------------------------------------------|-------------|--------------|----------------|--------------|---------------|--------------|--------------|---------------|
| (a) GPT-3 Confidence Intervals for Intersectional Differences - Probabilities of "No" (i.e. denying pain treatment) |             |              |                |              |               |              |              |               |
| 95% CI                                                                                                              | Asian Woman | Black Woman  | Hispanic Woman | White Woman  | Asian Man     | Black Man    | Hispanic Man | White Man     |
| Asian Woman                                                                                                         |             | -0.9% / 0.1% | -0.7% / 0.3%   | 0.4% / 1.2%  | 0.2% / 1.1%   | 0.3% / 1.2%  | 0.6% / 1.8%  | 0.9% / 1.9%   |
| Black Woman                                                                                                         |             |              | -0.3% / 0.6%   | 0.7% / 1.7%  | 0.6% / 1.5%   | 0.6% / 1.6%  | 1.0% / 2.1%  | 1.2% / 2.4%   |
| Hispanic Woman                                                                                                      |             |              |                | 0.5% / 1.6%  | 0.4% / 1.4%   | 0.5% / 1.4%  | 0.8% / 2.0%  | 1.0% / 2.3%   |
| White Woman                                                                                                         |             |              |                |              | -0.7% / 0.3%  | -0.6% / 0.4% | -0.2% / 1.0% | 0.1% / 1.0%   |
| Asian Man                                                                                                           |             |              |                |              |               | -0.4% / 0.6% | 0.0% / 1.1%  | 0.3% / 1.2%   |
| Black Man                                                                                                           |             |              |                |              |               |              | -0.1% / 1.0% | 0.1% / 1.2%   |
| Hispanic Man                                                                                                        |             |              |                |              |               |              |              | -0.4% / 0.8%  |
| Bonferroni-Corrected CI                                                                                             | Asian Woman | Black Woman  | Hispanic Woman | White Woman  | Asian Man     | Black Man    | Hispanic Man | White Man     |
| Asian Woman                                                                                                         |             | -1.2% / 0.4% | -1.0% / 0.6%   | 0.2% / 1.5%  | 0.0% / 1.4%   | 0.0% / 1.6%  | 0.2% / 2.2%  | 0.6% / 2.3%   |
| Black Woman                                                                                                         |             |              | -0.6% / 0.9%   | 0.4% / 2.1%  | 0.2% / 1.9%   | 0.3% / 1.9%  | 0.7% / 2.5%  | 0.8% / 2.8%   |
| Hispanic Woman                                                                                                      |             |              |                | 0.1% / 2.0%  | 0.0% / 1.7%   | 0.2% / 1.7%  | 0.4% / 2.4%  | 0.6% / 2.7%   |
| White Woman                                                                                                         |             |              |                |              | -1.0% / 0.7%  | -0.9% / 0.7% | -0.6% / 1.3% | -0.2% / 1.3%  |
| Asian Man                                                                                                           |             |              |                |              |               | -0.8% / 1.0% | -0.3% / 1.4% | 0.0% / 1.5%   |
| Black Man                                                                                                           |             |              |                |              |               |              | -0.5% / 1.4% | -0.3% / 1.6%  |
| Hispanic Man                                                                                                        |             |              |                |              |               |              |              | -0.7% / 1.1%  |
| (b) GPT-2 Confidence Intervals for Intersectional Differences - Probabilities of "No" (i.e. denying pain treatment) |             |              |                |              |               |              |              |               |
| 95% CI                                                                                                              | Asian Woman | Black Woman  | Hispanic Woman | White Woman  | Asian Man     | Black Man    | Hispanic Man | White Man     |
| Asian Woman                                                                                                         |             | -1.9% / 0.3% | -2.2% / 0.7%   | -1.3% / 1.1% | -2.7% / -0.5% | -0.4% / 2.4% | 0.2% / 2.8%  | -1.0% / 1.6%  |
| Black Woman                                                                                                         |             |              | -1.2% / 1.3%   | -0.2% / 1.7% | -1.5% / 0.0%  | 1.0% / 2.6%  | 1.3% / 3.3%  | 0.1% / 2.1%   |
| Hispanic Woman                                                                                                      |             |              |                | -0.5% / 1.9% | -2.0% / 0.4%  | 0.4% / 3.1%  | 1.0% / 3.5%  | -0.2% / 2.3%  |
| White Woman                                                                                                         |             |              |                |              | -2.2% / -0.7% | -0.2% / 2.3% | 0.7% / 2.5%  | -0.3% / 1.0%  |
| Asian Man                                                                                                           |             |              |                |              |               | 1.4% / 3.6%  | 2.3% / 3.8%  | 1.0% / 2.6%   |
| Black Man                                                                                                           |             |              |                |              |               |              | -0.7% / 1.8% | -1.9% / 0.5%  |
| Hispanic Man                                                                                                        |             |              |                |              |               |              |              | -2.1% / -0.3% |
| Bonferroni-Corrected CI                                                                                             | Asian Woman | Black Woman  | Hispanic Woman | White Woman  | Asian Man     | Black Man    | Hispanic Man | White Man     |
| Asian Woman                                                                                                         |             | -2.6% / 1.0% | -3.2% / 1.7%   | -2.1% / 1.9% | -3.4% / 0.3%  | -1.3% / 3.3% | -0.6% / 3.6% | -1.8% / 2.4%  |
| Black Woman                                                                                                         |             |              | -1.9% / 2.0%   | -0.8% / 2.3% | -1.9% / 0.5%  | 0.4% / 3.1%  | 0.7% / 3.9%  | -0.5% / 2.7%  |
| Hispanic Woman                                                                                                      |             |              |                | -1.3% / 2.7% | -2.7% / 1.1%  | -0.5% / 4.0% | 0.2% / 4.4%  | -1.0% / 3.1%  |
| White Woman                                                                                                         |             |              |                |              | -2.7% / -0.2% | -1.0% / 3.1% | 0.1% / 3.1%  | -0.7% / 1.5%  |
| Asian Man                                                                                                           |             |              |                |              |               | 0.7% / 4.3%  | 1.8% / 4.3%  | 0.5% / 3.2%   |
| Black Man                                                                                                           |             |              |                |              |               |              | -1.5% / 2.6% | -2.6% / 1.3%  |
| Hispanic Man                                                                                                        |             |              |                |              |               |              |              | -2.7% / 0.3%  |
| CI includes 0.e. is. not significant   CI Absolute Lower Bound > 0.5%   CI Absolute Lower Bound > 1.0%              |             |              |                |              |               |              |              |               |

CI includes 0% i.e. is not significant | CI Absolute Lower Bound > 0.5% | CI Absolute Lower Bound > 1.0%

Figure 3: This figure looks at *Group A v Group B* comparisons, posing the question: *is Group A at a disadvantage compared to Group B?* As such, a positive result implies that Group A is at a disadvantage compared to Group B, while a negative result implies the opposite. An interval including both negative and positive values is inconclusive. (a) 95% and Bonferroni-corrected Confidence Intervals (CI) for difference in probability of being denied pain treatment by GPT-3, (b) Similar for GPT-2.



### 4.3 Qualitative Assessment

We performed a brief qualitative assessment on the explanations produced by GPT-3 and GPT-2 to examine their consistency and medical validity (see Fig. 1c). Having access to a medical expert, we chose to evaluate the quality of the generated explanations according to (1) recognition of the correct pain diagnosis, (2) appropriate assessment of pain context, and (3) mention of appropriateness of opioid therapy.

We found that overall, GPT-3 generated well-constructed explanations, with no clear variation that could be ascribed to differences in gender or race, while GPT-2 was prone to incomplete and incorrect justifications, along with bits that merely parroted parts of the open prompt. It should be noted, however, that in several instances, both GPT-3 and GPT-2 failed to assign the correct name to the patient (e.g. replacing it with "Patient C") — this occurred mostly with Asian names.

## 5 Discussion

**Impact & Relevance:** The medical domain has had its fair share of AI mishaps, from algorithms exhibiting racial bias during risk prediction [34] to IBM’s Watson model making incorrect and unsafe treatment recommendations for cancer patients [35]. Yet, as AI becomes more advanced, many physicians are relying more heavily on intelligent models and algorithms to identify diseases and make medical decisions [36, 37]. In particular, LLMs are becoming so prominent that the UK’s government even considered replacing its national health emergency helpline with intelligent chatbots to provide medical advice to prospective patients [38].

Our results unfortunately align with well-known race and gender disparities around pain management [7] and confirm that using LLMs in medical contexts could reinforce detrimental conscious and unconscious biases still present in the healthcare system to this day. Furthermore, the fact that GPT-3 displayed no clear progress in the equity of its treatment of different demographic groups over GPT-2 is at the very least noteworthy, and evokes questions posed by previous researchers regarding the multifaceted costs of scaling LLMs [39]. There is a tendency for ethical AI work to be done retrospectively; with this study and the release of the Q-Pain dataset, we hope to arm researchers with the tools to further understand how bias can surface in medical QA systems before they are (possibly imminently) further deployed in the real world.

**Originality & Novelty:** The Q-Pain dataset can be used for experiments both in real-life medical settings (e.g. through surveys of clinicians) and with LLMs, thus allowing comparisons in significance and magnitude between the unconscious biases present in the healthcare system today and those found in language modeling. Additionally, to the best of our knowledge, this benchmark of GPT-3 and GPT-2—complete with rigorous statistical analysis—is the first of its kind in the context of medical treatment recommendations.

**Ethical Considerations:** The Q-Pain dataset is made of hypothetical scenarios and does not include any real patient’s personally identifiable information. While there is a risk related to publishing a dataset with limited scope—as it could be used counter to the paper’s intentions as proof that a model is conclusively not biased—we believe Q-Pain’s potential to help the AI community understand their progress in medical fairness domains outweighs the risk.

## 6 Limitations & Future Work

**Dataset Expansion:** Simply having a larger volume (approx. 500) of vignettes would increase the confidence in any results generated with this dataset. Additionally, we would like to provide greater diversity between medical scenarios, both within pain management and beyond, expanding the relevancy of our framework to address bias in medical decision-making.

**Closed Prompts:** In this setup, our goal was to standardize and neutralize the closed prompts to the greatest extent possible so as to isolate possible biases that arise without being primed to do so. But it might also be worthwhile to investigate how a QA system behaves when "provoked:" we would hope that even in the face of biased decision-making prompts, a QA system assisting in pain management decisions would not, itself, discriminate.

**Names:** The names we selected were derived using real-world data on demographic representations of first names [31], however demographic representation does not necessarily correlate with implicit stereotypical associations. Using a list of names more similar to that of the original implicit association test [30] could convey stronger stereotypical racial signals to the QA system, potentially yielding interesting results.

**Beyond Race & Gender:** Finally, it is worth mentioning our treatment of gender as binary and of race as four exhaustive and mutually exclusive groups: while it certainly made our setup more straightforward, in doing so, we may have missed other areas of treatment disparities and discrimination. Additionally, other social attributes related to culture, ethnicity, socioeconomic class, or sexual orientation have been shown to correlate with having poorer healthcare experiences [40] and would be worth exploring. What’s more, future work could involve creating a neutral baseline by not specifying either race or gender, or specifying a descriptor that isn’t a race or gender (e.g. "green," or "short"). This could add additional texture to the meaningfulness of our results and address potential concerns about the consistency of models across prompts.

## 7 Conclusion

In this study, we introduce Q-Pain, a dataset of 55 clinical vignettes designed specifically for assessing racial and gender bias in medical QA—whether in real-life medical settings or through AI systems—in the context of pain management, and offer an experimental framework using our dataset to target that research goal. We demonstrate the use of both dataset and framework on two LLMs, GPT-3 and GPT-2, and find statistically significant differences in treatment between intersectional race-gender subgroups, thus reaffirming the risk of using AI in medical settings.

Though ethical AI efforts are often performed retrospectively, with the release of the Q-Pain dataset, we hope to arm researchers with the tools to further understand how bias can surface in medical QA systems before they are further deployed in the real world.

## 8 Acknowledgements

We would like to thank Brad Ross for lending his statistical expertise to aid our analysis. We would also like to acknowledge Diana Zhao, MD, PhD and Jordan Atkins MD, MPH for their help in validating the dataset. Finally, we would like to thank Percy Liang, PhD from Stanford University for his support.

## References

- [1] Timothy H Wideman, Robert R Edwards, David M Walton, Marc O Martel, Anne Hudon, and David A Seminowicz. The multimodal assessment model of pain: a novel framework for further integrating the subjective pain experience within research and practice. *The Clinical journal of pain*, 35(3):212, 2019.
- [2] Jana M Mossey. Defining racial and ethnic disparities in pain management. *Clinical Orthopaedics and Related Research*, 469(7):1859–1870, 2011.
- [3] Robert C Coghill. Individual differences in the subjective experience of pain: new insights into mechanisms and models. *Headache: The Journal of Head and Face Pain*, 50(9):1531–1535, 2010.
- [4] Richard A Mularski, Foy White-Chu, Devorah Overbay, Lois Miller, Steven M Asch, and Linda Ganzini. Measuring pain as the 5th vital sign does not improve quality of pain management. *Journal of general internal medicine*, 21(6):607–612, 2006.
- [5] Rita M Holl and Jennifer Carmack. Complexity of pain, nurses’ knowledge, and treatment options. *Holistic nursing practice*, 29(6):377–380, 2015.
- [6] Brian B. Drwecki, Colleen F. Moore, Sandra E. Ward, and Kenneth M. Prkachin. Reducing racial disparities in pain treatment: The role of empathy and perspective-taking. *PAIN*, 152(5):1001–1006, 2011.
- [7] Karen O. Anderson, Carmen R. Green, and Richard Payne. Racial and ethnic disparities in pain: Causes and consequences of unequal care. *The Journal of Pain*, 10(12):1187–1204, 2009.

- [8] Carmen R. Green, Karen O. Anderson, Tamara A. Baker, Lisa C. Campbell, Sheila Decker, Roger B. Fillingim, Donna A. Kaloukalani, Kathryn E. Lasch, Cynthia Myers, Raymond C. Tait, Knox H. Todd, and April H. Vallerand. The Unequal Burden of Pain: Confronting Racial and Ethnic Disparities in Pain. *Pain Medicine*, 4(3):277–294, 09 2003.
- [9] Kelly M. Hoffman, Sophie Trawalter, Jordan R. Axt, and M. Norman Oliver. Racial bias in pain assessment and treatment recommendations, and false beliefs about biological differences between blacks and whites. *Proceedings of the National Academy of Sciences*, 113(16):4296–4301, 2016.
- [10] H. Majedi, S. Dehghani, Saeed Soleyman-jahi, A. Tafakhori, S. Emami, M. Mireskandari, and S. M. Hosseini. Assessment of factors predicting inadequate pain management in chronic pain patients. *Anesthesiology and Pain Medicine*, 9, 2019.
- [11] Anke Samulowitz, Ida Gremyr, Erik Eriksson, and Gunnel Hensing. “brave men” and “emotional women”: A theory-guided literature review on gender bias in health care and gendered norms towards patients with chronic pain. *Pain Research and Management*, 2018:1–14, 2018.
- [12] Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. BERT: Pre-training of deep bidirectional transformers for language understanding. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*, pages 4171–4186, Minneapolis, Minnesota, June 2019. Association for Computational Linguistics.
- [13] Tao Li, Daniel Khashabi, Tushar Khot, Ashish Sabharwal, and Vivek Srikumar. UNQOVERing stereotyping biases via underspecified questions. In *Findings of the Association for Computational Linguistics: EMNLP 2020*, pages 3475–3489, Online, November 2020. Association for Computational Linguistics.
- [14] Kris McGuffie and Alex Newhouse. The radicalization risks of GPT-3 and advanced neural language models. *CoRR*, abs/2009.06807, 2020.
- [15] Emily Sheng, Kai-Wei Chang, Premkumar Natarajan, and Nanyun Peng. The woman worked as a babysitter: On biases in language generation. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pages 3407–3412, Hong Kong, China, November 2019. Association for Computational Linguistics.
- [16] Emily Sheng, Kai-Wei Chang, Prem Natarajan, and Nanyun Peng. Towards Controllable Biases in Language Generation. In *Findings of the Association for Computational Linguistics: EMNLP 2020*, pages 3239–3254, Online, November 2020. Association for Computational Linguistics.
- [17] Nikita Nangia, Clara Vania, Rasika Bhalerao, and Samuel R. Bowman. CrowS-pairs: A challenge dataset for measuring social biases in masked language models. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 1953–1967, Online, November 2020. Association for Computational Linguistics.
- [18] Moin Nadeem, Anna Bethke, and Siva Reddy. Stereoset: Measuring stereotypical bias in pretrained language models, 2020.
- [19] Alec Radford, Jeff Wu, Rewon Child, David Luan, Dario Amodei, and Ilya Sutskever. Language models are unsupervised multitask learners. 2019.
- [20] Tom Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared D Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, Sandhini Agarwal, Ariel Herbert-Voss, Gretchen Krueger, Tom Henighan, Rewon Child, Aditya Ramesh, Daniel Ziegler, Jeffrey Wu, Clemens Winter, Chris Hesse, Mark Chen, Eric Sigler, Mateusz Litwin, Scott Gray, Benjamin Chess, Jack Clark, Christopher Berner, Sam McCandlish, Alec Radford, Ilya Sutskever, and Dario Amodei. Language models are few-shot learners. In H. Larochelle, M. Ranzato, R. Hadsell, M. F. Balcan, and H. Lin, editors, *Advances in Neural Information Processing Systems*, volume 33, pages 1877–1901. Curran Associates, Inc., 2020.
- [21] David Vilares and Carlos Gómez-Rodríguez. HEAD-QA: A healthcare dataset for complex reasoning. In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, pages 960–966, Florence, Italy, July 2019. Association for Computational Linguistics.
- [22] Asma Ben Abacha and Dina Demner-Fushman. A question-entailment approach to question answering. *BMC bioinformatics*, 20(1):1–23, 2019.

- [23] Asma Ben Abacha, Yassine Mrabet, Mark Sharp, Travis Goodwin, Sonya E. Shooshan, and Dina Demner-Fushman. Bridging the gap between consumers' medication questions and trusted answers. In *MEDINFO 2019*, 2019.
- [24] Anusri Pampari, Preethi Raghavan, Jennifer Liang, and Jian Peng. emrQA: A large corpus for question answering on electronic medical records. In *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing*, pages 2357–2368, Brussels, Belgium, October–November 2018. Association for Computational Linguistics.
- [25] Xiao Zhang, Ji Wu, Zhiyang He, Xien Liu, and Ying Su. Medical exam question answering with large-scale reading comprehension. 2018.
- [26] Paulyne Lee, Maxine Le Saux, Rebecca Siegel, Monika Goyal, Chen Chen, Yan Ma, and Andrew C. Meltzer. Racial and ethnic disparities in the management of acute pain in us emergency departments: Meta-analysis and systematic review. *The American Journal of Emergency Medicine*, 37(9):1770–1777, 2019.
- [27] Carol S Weisse, Paul C Sorum, Kafi N Sanders, and Beth L Syat. Do gender and race affect decisions about pain management? *Journal of general internal medicine*, 16(4):211–217, 2001.
- [28] Haoran Zhang, Amy X. Lu, Mohamed Abdalla, Matthew McDermott, and Marzyeh Ghassemi. Hurtful words: Quantifying biases in clinical contextual word embeddings. In *Proceedings of the ACM Conference on Health, Inference, and Learning*, CHIL '20, page 110–120, New York, NY, USA, 2020. Association for Computing Machinery.
- [29] A. L. Goldberger, L. A. N. Amaral, L. Glass, J. M. Hausdorff, P. Ch. Ivanov, R. G. Mark, J. E. Mietus, G. B. Moody, C.-K. Peng, and H. E. Stanley. PhysioBank, PhysioToolkit, and PhysioNet: Components of a new research resource for complex physiologic signals. *Circulation*, 101(23):e215–e220, 2000. Circulation Electronic Pages: <http://circ.ahajournals.org/content/101/23/e215.full> PMID:1085218; doi: 10.1161/01.CIR.101.23.e215.
- [30] Anthony G Greenwald, Debbie E McGhee, and Jordan LK Schwartz. Measuring individual differences in implicit cognition: the implicit association test. *Journal of personality and social psychology*, 74(6):1464, 1998.
- [31] Konstantinos Tzioumis. Demographic aspects of first names. *Scientific data*, 5(1):1–9, 2018.
- [32] Kimberle Crenshaw. Demarginalizing the intersection of race and sex: A black feminist critique of antidiscrimination doctrine, feminist theory and antiracist politics. *University of Chicago Legal Forum*, Vol. 1989 , Article 8, 1989.
- [33] Kawin Ethayarajh. Is your classifier actually biased? measuring fairness under uncertainty with bernstein bounds. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 2914–2919, Online, July 2020. Association for Computational Linguistics.
- [34] Jenna Wiens, W. Nicholson Price, and Michael W. Sjoding. Diagnosing bias in data-driven algorithms for healthcare. *Nature Medicine*, 26(1):25–26, 2020.
- [35] Eliza Strickland. Ibm watson, heal thyself: How ibm overpromised and underdelivered on ai health care. *IEEE Spectrum*, 56(4):24–31, 2019.
- [36] Michael D. Abràmoff, Philip T. Lavin, Michele Birch, Nilay Shah, and James C. Folk. Pivotal trial of an autonomous ai-based diagnostic system for detection of diabetic retinopathy in primary care offices. *npj Digital Medicine*, 1(1), 2018.
- [37] Seung Seog Han, Ilwoo Park, Sung Eun Chang, Woohyung Lim, Myoung Shin Kim, Gyeong Hun Park, Je Byeong Chae, Chang Hun Huh, and Jung-Im Na. Augmented intelligence dermatology: deep neural networks empower medical professionals in diagnosing skin cancer and predicting treatment options for 134 skin disorders. *Journal of Investigative Dermatology*, 140(9):1753–1761, 2020.
- [38] Madhumita Murgia. Nhs to trial artificial intelligence app in place of 111 helpline. *Financial Times*, 2017.
- [39] Emily M. Bender, Timnit Gebru, Angelina McMillan-Major, and Shmargaret Shmitchell. On the dangers of stochastic parrots: Can language models be too big? In Madeleine Clare Elish, William Isaac, and Richard S. Zemel, editors, *FAccT '21: 2021 ACM Conference on Fairness, Accountability, and Transparency, Virtual Event / Toronto, Canada, March 3-10, 2021*, pages 610–623. ACM, 2021.
- [40] Kenneth D. Craig, Cindy Holmes, Maria Hudspeth, Gregg Moor, Mehmoona Moosa-Mitha, Colleen Varcoc, and Bruce Wallace. Pain in persons who are marginalized by social conditions. *Pain*, 161(2):261–265, 2019.